



Assessment Report
on
“Predict Product Return”
submitted as partial fulfillment for the award of
BACHELOR OF TECHNOLOGY
DEGREE

SESSION 2024-25

in
CSE(AI)

By

Name : Nirmal Kumar

Roll Number : 202401100300162

Section: C

Under the supervision of
“MAYANK LAKHOTIYA”

KIET Group of Institutions, Ghaziabad

April, 2025

Introduction

In the modern retail landscape, customer satisfaction and supply chain efficiency play a crucial role in driving business success. One key challenge many businesses face is the high volume of **product returns**, which can disrupt inventory, increase costs, and impact customer experience. Understanding the **patterns behind product returns** is essential for improving logistics, enhancing product offerings, and refining marketing strategies.

This project focuses on using **machine learning and data analysis techniques** to explore, visualize, and model customer return behavior based on historical purchase data. It leverages both **classification models** (to predict return likelihood) and **clustering algorithms** (to segment customers or products when return labels are unavailable).

Key objectives of this analysis include:

- Identifying significant factors that influence product returns.
- Evaluating a logistic regression model using metrics such as **accuracy, precision, recall**, and **AUC**.
- Visualizing feature distributions, correlations, and confusion matrices.
- Segmenting customer/product behavior using unsupervised **K-Means clustering** when the target variable is missing.

The outcome of this project helps stakeholders make informed decisions regarding **product design, delivery optimization, and customer service strategies** based on actionable insights derived from data.

Methodology

The methodology of this project is divided into two key branches based on the nature of the data: **classification** (when return labels are available) and **unsupervised clustering** (when return labels are missing). Below is a step-by-step breakdown of the entire workflow:

1. Data Loading & Preprocessing

- The dataset (`product_return.csv`) is loaded using **Pandas**.
- Initial inspection includes checking for **missing values**, data types, and previewing the structure of the dataset.
- The `returned` column (Yes/No) is **encoded into binary values** (1 for return, 0 for no return).
- Numerical features are standardized using **StandardScaler** for better model performance.

2. Exploratory Data Analysis (EDA)

- **Visualizations** such as histograms, box plots, and count plots are generated to understand feature distributions and potential outliers.
- A **correlation heatmap** is plotted to identify relationships between numerical features and the target variable (`returned`).
- Feature importance is later analyzed through logistic regression coefficients.

3. Classification (Supervised Learning)

- The dataset is split into **training and test sets** (80/20 split).
- A **Logistic Regression** model is trained on the scaled training data.
- **Model predictions** are made on the test set.
- Evaluation metrics calculated:
 - **Accuracy**

- **Precision**
- **Recall**
- **Classification Report**
- **Confusion Matrix** (both raw and normalized)
- **ROC Curve & AUC**
- **Precision-Recall Curve**

4. Clustering (Unsupervised Learning)

- When return labels are not available, the model switches to **K-Means clustering**.
- Numerical features are extracted and scaled.
- K-Means is applied to segment the data into 3 clusters (`n_clusters=3`).
- The resulting clusters are visualized using **pair plots**, and the **cluster centroids** are analyzed to interpret typical segment characteristics.

5. Visualization & Interpretation

- Extensive use of **Seaborn** and **Matplotlib** to visualize:
 - Feature distributions by class
 - Confusion matrix heatmaps
 - Model evaluation curves
 - Feature importance rankings
 - Clustering groupings

This approach ensures a comprehensive analysis of product return behavior, adaptable to both labeled and unlabeled datasets, while delivering clear visual and statistical insights for further business decision-making.

CODE TYPED

```
# Import Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    classification_report, confusion_matrix, accuracy_score,
    precision_score, recall_score, roc_curve, auc, precision_recall_curve
)
from sklearn.cluster import KMeans

# Set aesthetics
sns.set(style='whitegrid')
plt.rcParams['figure.figsize'] = (8, 5)

# Load Dataset
df = pd.read_csv('product_return.csv')
print("First 5 rows:\n", df.head())

# Detect task type: classification or clustering
task_type = 'classification' if 'returned' in df.columns else 'clustering'

if task_type == 'classification':
    print("\nDetected Classification Task")

    # Check for missing values
    print("\nMissing Values:\n", df.isnull().sum())

    # Encode target
    df['returned'] = LabelEncoder().fit_transform(df['returned']) #
yes=1, no=0

    # Summary statistics
    print("\nDescriptive Statistics:\n", df.describe())
```

```

# Target Distribution
sns.countplot(data=df, x='returned')
plt.title("Distribution of Returned Products (0 = No, 1 = Yes)")
plt.show()

# Histograms
for col in ['purchase_amount', 'review_score', 'days_to_delivery']:
    sns.histplot(data=df, x=col, hue='returned', kde=True, bins=20,
palette='Set2')
    plt.title(f'Distribution of {col} by Return Status')
    plt.show()

# Boxplots
for col in ['purchase_amount', 'review_score', 'days_to_delivery']:
    sns.boxplot(data=df, x='returned', y=col, palette='Set3')
    plt.title(f'Boxplot of {col} by Return Status')
    plt.show()

# Correlation Heatmap
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

# Prepare Data
X = df.drop('returned', axis=1)
y = df['returned']

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Feature Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train Model
model = LogisticRegression()
model.fit(X_train_scaled, y_train)

```

```

# Predictions and Probabilities
y_pred = model.predict(X_test_scaled)
y_proba = model.predict_proba(X_test_scaled)[:, 1]

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
cm_normalized = confusion_matrix(y_test, y_pred, normalize='true')

fig, ax = plt.subplots(1, 2, figsize=(12, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax[0],
            xticklabels=['No Return', 'Return'],
            yticklabels=['No Return', 'Return'])
ax[0].set_title('Confusion Matrix (Raw)')
ax[0].set_xlabel('Predicted')
ax[0].set_ylabel('Actual')

sns.heatmap(cm_normalized, annot=True, fmt='.2f', cmap='Greens',
ax=ax[1],
            xticklabels=['No Return', 'Return'],
            yticklabels=['No Return', 'Return'])
ax[1].set_title('Confusion Matrix (Normalized)')
ax[1].set_xlabel('Predicted')
ax[1].set_ylabel('Actual')

plt.tight_layout()
plt.show()

# Metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)

print(f"\nAccuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
print("\nClassification Report:\n", classification_report(y_test,
y_pred, target_names=['No Return', 'Return']))

# ROC Curve

```

```

fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = auc(fpr, tpr)
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.grid(True)
plt.show()

# Precision-Recall Curve
precision_vals, recall_vals, _ = precision_recall_curve(y_test,
y_proba)
plt.plot(recall_vals, precision_vals, marker='.')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.grid(True)
plt.show()

# Feature Importance
coef = model.coef_[0]
feature_df = pd.DataFrame({'Feature': X.columns, 'Importance': coef})
feature_df = feature_df.sort_values(by='Importance', key=abs,
ascending=False)
sns.barplot(data=feature_df, x='Importance', y='Feature',
palette='viridis')
plt.title('Feature Importance (Logistic Coefficients)')
plt.axvline(0, color='black', lw=1)
plt.show()

else:

    print("\nDetected Clustering Task")

    # Choose numeric features
    X = df.select_dtypes(include='number')

    # Scale features
    scaler = StandardScaler()

```



```
X_scaled = scaler.fit_transform(X)

# KMeans Clustering
kmeans = KMeans(n_clusters=3, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)

# Visualize Clusters
sns.pairplot(df, hue='Cluster', palette='tab10')
plt.suptitle('KMeans Clusters Visualization', y=1.02)
plt.show()

# Cluster Centers
centers =
pd.DataFrame(scaler.inverse_transform(kmeans.cluster_centers_),
columns=X.columns)
print("\nCluster Centers:\n", centers)
```

OUTPUT

```
(4) First 5 rows:
   purchase_amount  review_score  days_to_delivery  returned
0      687.011818      3.778615           4         no
1      325.972093      2.458683           1         yes
2      685.382724      3.954024           7         no
3      291.100577      3.666468          14         yes
4      209.806672      1.478248           2         no
```




Detected Classification Task


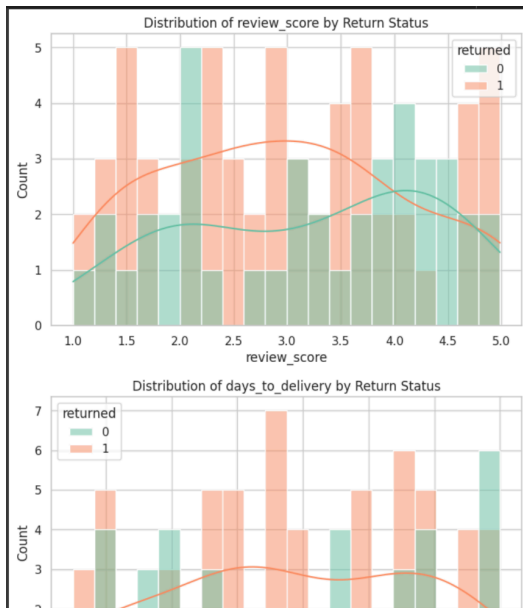
Missing Values:


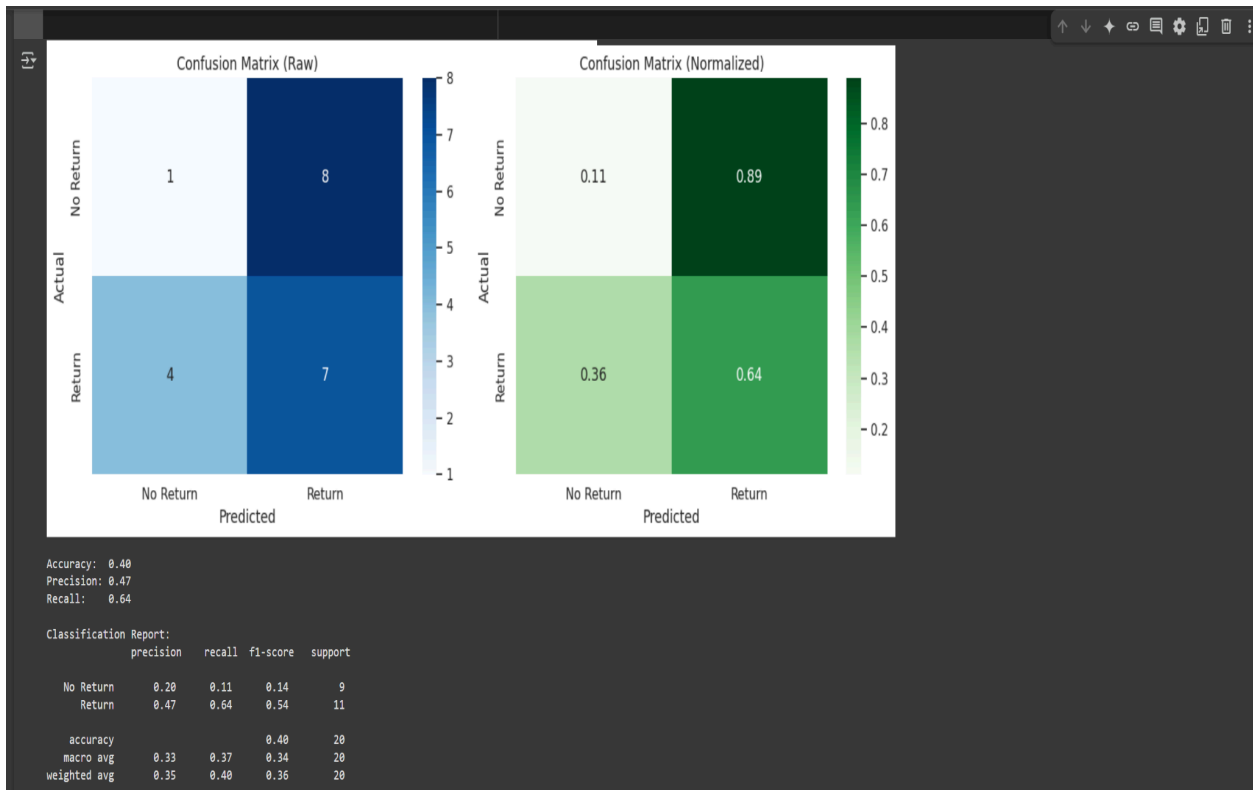
```
purchase_amount    0
review_score        0
days_to_delivery   0
returned            0
dtype: int64
```


Descriptive Statistics:

	purchase_amount	review_score	days_to_delivery	returned
count	100.000000	100.000000	100.000000	100.000000
mean	573.396104	3.065745	7.830000	0.590000
std	267.802559	1.160948	4.119564	0.494311
min	112.270622	1.000746	1.000000	0.000000
25%	351.637815	2.124633	4.000000	0.000000
50%	598.360894	3.097644	7.500000	1.000000
75%	827.319433	4.072944	11.250000	1.000000
max	997.218297	4.986263	14.000000	1.000000










<ipython-input-9-71dfa0ef0f1d>:136: FutureWarning:

Passing 'palette' without assigning 'hue' is deprecated and will be removed in v0.14.0. Assign the 'y' variable to 'hue' and set 'legend=False' for the same effect.

```
sns.barplot(data=feature_df, x='Importance', y='Feature', palette='viridis')
```



2s completed at 11:13AM

References

1. **Scikit-learn Documentation**
<https://scikit-learn.org/stable/documentation.html>
Used for model training, evaluation metrics, preprocessing, and clustering algorithms.
2. **Pandas Documentation**
<https://pandas.pydata.org/docs/>
Used for data manipulation and preprocessing.
3. **Seaborn Documentation**
<https://seaborn.pydata.org/>
Used for statistical data visualization.
4. **Matplotlib Documentation**
<https://matplotlib.org/stable/contents.html>
Used for plotting graphs and model evaluation curves.
5. **Géron, A. (2019). *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow* (2nd ed.). O'Reilly Media.**
Conceptual guidance on classification and clustering in Python.
6. **Tan, P.-N., Steinbach, M., & Kumar, V. (2018). *Introduction to Data Mining* (2nd ed.). Pearson.**
Understanding clustering techniques and unsupervised learning fundamentals.
7. **IBM Cloud Education: What is a Confusion Matrix?**
Overview of confusion matrices and model evaluation.